

Report for CMU Data Science Cup 2016

Team The Third Place

I. FINDINGS

II. DATA EXPLORATION

We took multiple steps to first clean and process our data. Firstly, we noticed that there were multiple entries in which the QUANTITY, as well as other relevant variables like the BASE_SPEND_AMT and NET_SPEND_AMT were all 0. This made up approximately 3000 of the 300,000 total rows of our dataset. Furthermore, we normalized fields like the DAY from 500 to 700 to 0 to 20. This makes our data more palatable to us and the models. Finally, we added a boolean field, "GET_EGGS", that represented whether or not the current entry was an egg purchase or not.

Then, we took each grocery transaction, and aggregated them together into tables representing transactions by day, and then transactions by month, per household. This was done by taking a particular day/month, aggregating the QUANTITY and SPEND fields, and appending metadata for the purchases.

For feature engineering, we explored the relationships between different pairs of variables, in order to decide what features would be relevant to the task. Examples:

III. MODEL EXPERIMENTS

In this competition, we explored the results of three strong state-of-the-art methods in classification: logistic regression, random forests, and Gradient Boosting Machines. We trained each model with the same set of features:

IV. BASELINE

We have two baseline models used for comparison. The first one is based off the observation that our output parameter of the data is heavily skewed.

The second baseline is trying the probability that each household purchasing eggs as a bernoulli random variable, so we have a binomial distribution in this case. The conditional probability is given by:

$$P\{GET\ EGGS\} = \frac{\sum_i^{household} \sum_j^{day} GET\ EGGS_{ij}}{\sum_i^{household} \sum_j^{day} 1}$$

$$P\{GET\ EGGS\} | week = 1 - P\{GETNOEGG\}$$

$$= 1 - (1 - P\{GET\ EGGS\})^7$$

$$= 0.126284$$
